

FIT5037: Network Security

Security at Application layer

Faculty of Information Technology
Monash University

April 30, 2019

Commonwealth of Australia (*Copyright Regulations 1969*)

Warning: This material has been reproduced and communicated to you by or on behalf of Monash University pursuant to Part VB of the *Copyright Act 1968 (the Act)*. The material in this communication may be subject to copyright under the Act. Any further reproduction of communication of the material by you may be subject of copyright protection under the Act.

Do not remove this notice.

Lecture 8: Security at Application layer

Lecture Topics:

- Symmetric key cryptography
- Asymmetric key cryptography
- Pseudorandom Number Generators and hash functions
- Authentication Methods and AAA protocols
- Security at Network layer (IPsec)
- Security at Network layer (firewalls and wireless security)
- Security at Transport layer
- **Security at Application layer**
- Computer system security and malicious code
- Computer system vulnerabilities and penetration testing
- Intrusion detection
- Denial of Service Attacks and Countermeasures / Revision

Outline

- Electronic Mail Security
 - brief review of PGP and S/MIME
 - mention of DKIM
- Secure Shell Protocol
- DNS Service Security

Security of Electronic Mail Messages

Electronic Mail Security

- The basis for email over the Internet
 - Simple Mail Transfer Protocol (SMTP originally specified in RFC821 August 1982, latest: RFC5321 October 2008)
 - Message syntax (originally specified in RFC822, latest: 5322)
 - Multipurpose Internet Mail Extension (MIME specified in RFC 2045-2049)
- Neither SMTP nor the message syntax supports security services
- Required security properties:
 - confidentiality: protection from disclosure
 - authentication: of sender of message
 - message integrity: protection from modification
 - non-repudiation of origin: protection from denial by sender
- Why providing transport layer security is not enough for SMTP protocol?
- Two main schemes for end-to-end security of email messages:
 - PGP
 - S/MIME


Pretty Good Privacy (PGP)

- PGP was developed by Philip R. Zimmerman
- PGP provides confidentiality and authentication services that can be used for electronic mail and file storage applications
- RFC 4880: OpenPGP Message Format is an Internet standard
 - Referred to as PGP 5.x
 - GnuPG or GPG is an OpenPGP implementation (available as an open source project)
- PGP General Functions
 - key management
 - key ring for public and private key storage
 - per-message session keys
 - authentication and data integrity
 - digital signature
 - compression performed
 - after digital signature
 - before encryption
 - confidentiality
 - encryption
 - Radix-64 (Base64) conversion
 - interoperability of text representation between different mail servers and relay agents

PGP: Authentication and Confidentiality (Sender)

- PGP can provide authentication only or authentication and confidentiality


Alice: Sender



Providing both services on the same message:

- Sender first signs the message with its own private key
- Encrypts compressed (message + signature) with a session key
- Encrypts the session key with the recipient's public key
- and the encrypted key precedes the rest

PGP: Authentication and Confidentiality (Receiver)



- Receiver first decrypts the session key using his private key
- Decrypts compressed (message + signature) with the recovered session key
- Verifies the signature with sender's public key
- Stores message + signature if verified

PGP Public Key Management: Web of Trust

- Rather than relying on certificate authorities, in PGP every user is one's own CA
 - can sign keys for users they know directly
- Forms a “web of trust”
 - trusted keys will be signed
 - can trust keys others have signed if there is chain of signatures to them
- Key ring includes trust indicators
- Users can also revoke their keys

Secure Multipurpose Internet Mail Extensions (S/MIME)

- original SMTP protocol can only transmit text messages
- MIME provides support for varying content types and multi-part messages
 - encodes binary data in textual form
 - converts 8-bit binary values to printable characters
- S/MIME¹ added security enhancements to MIME
 - is very similar to PGP
- creates a MIME body according to Cryptographic Message Syntax (CMS)²
- Also offers the ability to sign and/or encrypt messages with the following functions
 - **enveloped data**: encrypted content and associated keys
 - **signed data**: encoded (message + signed digest)
 - **clear-signed data**: clear-text message + encoded signed digest
 - recipient with MIME capability but not S/MIME capability can read the message
 - **signed and enveloped data**: nesting of signed and encrypted entities, i.e., various orderings for encrypting and signing

¹“S/MIME version 3.2”

²Cryptographic Message Syntax

S/MIME Certificate Use

- S/MIME uses X.509 v3 certificates
- managed by using a hybrid of a strict X.509 CA hierarchy and PGP's web of trust
- each client has a list of trusted CA's certificates
- and own public/private key pairs and certificates
- certificates must be signed by trusted CA's
- S/MIME agents must provide some certificate retrieval mechanism
 - using X.500 directory service
 - using secure DNS to associate certificates with domain names for S/MIME (RFC8162)

Domain Keys Identified Mail (DKIM)

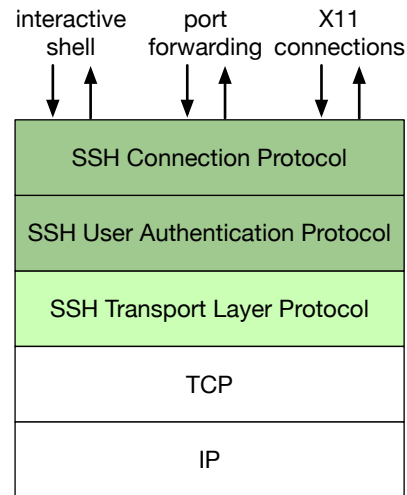
- a specification for cryptographically signing email messages by domains rather than users
- so signing domain claims responsibility
 - ADMD: Administrative Management Domain
 - SDMD: Signing Domain Identifier
- recipients / agents can verify signature
 - DKIM separates the identity of the signer from the author of the message
 - the recipient can interpret the header fields of DKIM signatures and decide how to process the message
- proposed Internet Standard (latest) RFC6376
 - has been widely adopted
 - provides authenticity and message integrity for the header fields as well as the body of the messages
 - relies on Domain Name Services (DNS) for key management
 - does not provide confidentiality

Secure Shell (SSH)

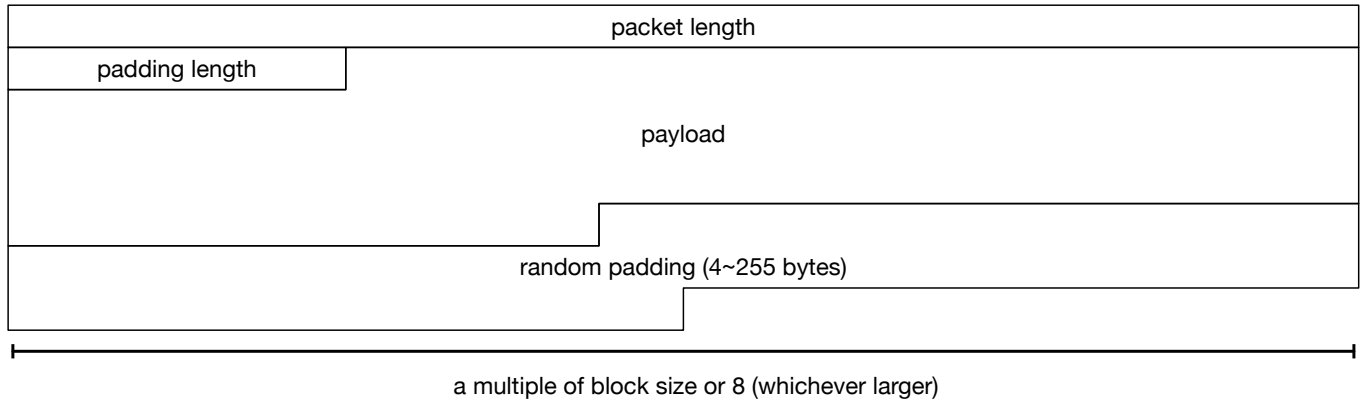
Secure Shell: A secure protocol for remote login

Secure Shell (SSH) Protocol

- protocol architecture is defined in RFC 4251
- runs over a reliable transport layer protocol (such as TCP)
- comprised of three major components:
 - Transport Layer Protocol defined in RFC 4253
 - server authentication
 - confidentiality
 - integrity
 - perfect forward secrecy
 - optional compression
 - (client) Authentication Protocol defined in RFC 4252
 - authenticates the client to the server
 - Connection Protocol defined in RFC 4254
 - multiplexes the encrypted tunnel into several logical channels



SSH Transport Layer Protocol: General Message Structure



SSH Transport Layer Protocol: Confidentiality and Integrity

$\text{message} = \text{packet_length} \parallel \text{padding_length} \parallel \text{payload} \parallel \text{padding}$

$C = E(K, \text{message})$

$T = \text{MAC}(K, \text{sequence_number} \parallel \text{message})$

- **sequence_number**: a 32-bit unsigned integer which is implicit (not included with the message)
 - set to 0 for the first message
 - incremented for every message regardless of encryption or MAC being used
- MAC algorithm is chosen independently at each side
 - recommended to be the same in practice
- MAC is transmitted unencrypted $C \parallel T$
 - this approach is referred to as Encrypt-and-MAC and is vulnerable for some combinations of ciphers and MACs³

³Bellare, Mihir, Tadayoshi Kohno, and Chanathip Namprempre. "Authenticated encryption in SSH: provably fixing the SSH binary packet protocol." Proceedings of the 9th ACM conference on Computer and communications security. ACM 2002.

SSH Transport Layer Protocol: Key Exchange

- Kex starts with each side sending list of supported algorithms
 - if the first algorithm is the same in both client and server list it will be selected
 - otherwise server will iterate through client list to find one that server also supports
 - if explicit server authentication is used the key exchange messages are signed
- each side may guess the kex algorithm and send its public key
 - `first_kex_packet_follows` indicates this
 - if the guess is wrong the packet will be ignored
- `kex_algorithms`: DH kex and a hash function
 - initially modulo prime DH groups supported
 - support for ECC public key and ECDH kex are added in RFC 5656
 - group parameters are defined for the protocol (separate from IKE groups)
 - hash function is used to derive keys from shared secret (Appendix)
- RFC 6239: Suite B Cryptographic Suites for Secure Shell (SSH)

client and server key exchange message:


byte	SSH_MSG_KEXINIT
byte[16]	cookie (random bytes)
name-list	kex_algorithms
name-list	server_host_key_algorithms
name-list	encryption_algorithms_client_to_server
name-list	encryption_algorithms_server_to_client
name-list	mac_algorithms_client_to_server
name-list	mac_algorithms_server_to_client
name-list	compression_algorithms_client_to_server
name-list	compression_algorithms_server_to_client
name-list	languages_client_to_server
name-list	languages_server_to_client
boolean	first_kex_packet_follows
uint32	0 (reserved for future extension)

SSH Transport Layer Protocol: Public Key Authentication

- supports public key method for server authentication
 - defines key format, signature and encryption algorithm, encoding of signature and encrypted data
 - DSS and RSA raw keys
 - OpenPGP certificates
 - X.509 certificates (RFC 6187)
- client may have received the public key of the server (host key) securely (outside protocol)
 - otherwise the protocol is vulnerable to MitM attack
- server signs its key exchange message with its private key
 - the key exchange also contains the host (public) key
- client verifies the host key
 - this authenticates the shared secret from kex
 - provided that public key of server is authentic

SSH Transport Layer Protocol: Kex and Auth

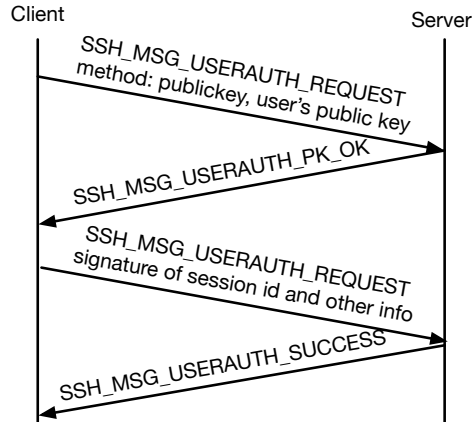
- In the following exchanges:
 - C : client:
 - S : server
 - V_C : client id
 - V_S : server id
 - PU_S : server's public (host) key
- SSH_MSG_KEXINIT negotiate supported algorithms
 - message structure shown in a previous slide
 - I_C : client's kex init message
 - I_S : server's kex init message
- SSH_MSG_KEXDH_INIT: DH_C
 - DH public key of the client
- SSH_MSG_KEXDH_REPLY: $PU_S || DH_S || sig$
 - $H = hash(V_C || V_S || I_C || I_S || PU_S || DH_C || DH_S || K)$
 - $sig = Sign(PR_S, H)$
 - DH_S : is server's DH public key
 - K is: the DH shared secret
- Encryption and MAC keys are derived from K (Appendix)



SSH (client) Authentication Protocol

- receives from the lower level:
 - the session identifier
 - whether confidentiality is provided
 - client starts with `SSH_MSG_USERAUTH_REQUEST`
 - providing a user name
 - service name (there may be more than one service provided)
 - method name
 - method specific fields
 - server tells client what authentication methods can be used
 - publickey
 - password
 - hostbased
 - none
 - additional methods can be defined
 - server responds with
 - `SSH_MSG_USERAUTH_FAILURE` if request is rejected
 - whether authentication can continue providing a list of methods
 - `SSH_MSG_USERAUTH_SUCCESS` if request is accepted
- is sent only after authentication sequence is completed

client publickey authentication



SSH Connection Protocol

- runs on top of SSH Transport Layer and (client) Authentication protocols
- multiplexes different types of user communication using channels over the same transport layer connection
 - terminal sessions, forwarded TCP connections, X11 forwarding etc.
- channels are identified by numbers
- each side can open a channel by sending `SSH_MSG_CHANNEL_OPEN`
 - `channel type`: a string specifying the type of channel
 - `sender channel`: local identifier of the channel
 - `initial window size`: how many bytes can be sent without adjusting windows (flow control)
 - `maximum packet size`
 - `channel data`
- other side responds with
 - `SSH_MSG_CHANNEL_OPEN_CONFIRMATION`
 - or `SSH_MSG_CHANNEL_OPEN_FAILURE`
- further channel data is sent with `SSH_MSG_CHANNEL_DATA` messages
- when no more data is sent by a party it will send `SSH_MSG_CHANNEL_EOF`
- when either party wishes to terminate a channel it sends `SSH_MSG_CHANNEL_CLOSE`

Domain Name Service

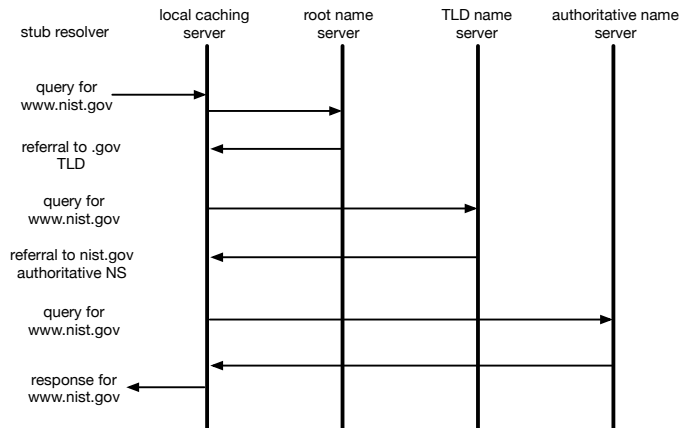
Security of Domain Name Service

Domain Name Service (DNS)

- DNS resolves domain names to network (IP) addresses
 - for network devices IP addresses are unique identifiers
 - to access resources using user-friendly names domain names are used
- DNS infrastructure is comprised of geographically distributed entities
 - uses a hierarchical model
 - *root domain* is represented as a dot ("."): root servers resolve the next level queries
 - *Top Level Domain (TLD)*: TLD servers resolve second-level domains
 - *authoritative* name servers for each domain resolve names under that particular domain
 - DNS cache servers store a copy of resolved names to provide a faster response
- DNS data is intended to be accessible to any device in the Internet
 - confidentiality for public information is not a concern
 - integrity and origin authenticity of DNS information is the main goal for DNS security

DNS Request Resolution Process

- DNS resolution process as described in NIST SP-800-81-2:



- *stub resolver*: a process resolving user queries
- *recursive name server* or *resolving name server*: a caching DNS server
 - will check its cache regarding requested name and responds if found
 - will recursively check the cache or query TLD/root servers until query is resolved
 - will send back an error if not found after all queries

Attacks on DNS

- Exploit DNS host
- Masquerade
 - an adversary spoofs the identity of a DNS node
 - denies access to services identified by the spoofed node service
 - provides bogus information poisoning DNS cache servers
- Violation of integrity of DNS information
 - stored on the authoritative DNS
 - stored on a DNS cache server
 - in transit
- Use DNS as a reflect/amplify Distributed Denial of Service attack
 - DNS uses UDP protocol vulnerable to this type of attack (spoofed source IP)
 - no connection is required
 - response will be sent to spoofed IP

DNS Security

- Secure DNS host
 - secure operating system
 - patch/update OS
 - least privilege etc.
 - secure applications
 - remove unnecessary software
 - patch/update software
 - process isolation
 - network fault tolerance
 - server redundancy
- Protect DNS management transactions
 - outlined as TSIG in RFC 2845 and updated by RFC 3645
 - update of DNS records
 - data replication between DNS nodes
- Protect DNS query/response transactions
 - outlined as DNSSEC in RFC 4033
 - can offload the cryptographic operations from clients to DNS servers
- NIST SP-800-81-2: Secure Domain Name System (DNS) Deployment Guide

- **Zone:** a configurable entity within a name server that contains the domain name information
 - Owner name: domain name
 - TTL: time to live
 - Class: currently one class IN
 - RRTYPE: Resource Record Types defined in RFC 1035 (and updated by several RFCs)
 - SOA: marks the start of a zone authority
 - A: alias, provides the IP address for a host name
 - MX: mail exchanger
 - NS: an authoritative name server
 - CNAME: a canonical name for an alias
 - RData: Resource information

DNS Server Types

- authoritative: server is an authoritative source for RRs of a particular zone
 - *master* or *primary*: contains (authoritative) zone files created and edited manually by zone administrator or updated dynamically by trusted clients
 - *slave* or *secondary*: also contains (authoritative) zone files created through replication from master server by *zone transfer* DNS transaction
 - added to provide fault-tolerance
 - zone transfer: is a name resolution query with type code AXFR meaning “all RRs for the zone”
 - the slave server is notified by *NOTIFY* transaction when a zone file in master server changes
- caching (also resolving/recursive): resolves name queries by asking authoritative name servers (in the hierarchy) or from its built-up cache (from previous queries)
 - generally local name server resolving queries on behalf of local clients
 - a name server can be both authoritative and caching
 - an authoritative DNS server that resolves recursive queries may be vulnerable to compromise of its zone information

DNS Transactions

- DNS query/response
 - a query is made by a stub resolver
 - a response may come from an authoritative or caching server
- Zone Transfer
 - the method by which the secondary server updates its zone files
 - a zone transfer query requests all RRs from the primary server
 - originates from the secondary server
 - either as a response to DNS NOTIFY
 - or based on the value of Refresh field in SOA section of zone file
 - this transaction reveals much more information than a DNS query
- Dynamic Updates
 - add/delete individual RRs, delete specific set of RRs e.g. same owner name, RRType etc., delete an existing domain (all RRs), add a new domain
 - when enabled open the server to various attacks
- DNS NOTIFY
 - signals the secondary server to initiate a zone transfer
 - more efficient than frequent polling
 - fraudulent DNS NOTIFY messages could trigger unnecessary zone transfer (may require transmission of large amount of information), can be prevented by limiting the sender to primary server

DNS Data Contents: Threats

- **Lame Delegation**
 - FQDN or IP address of name servers are changed in child zone but not in the parent zone
 - child zone becomes unreachable
- **Zone Drift**
 - Refresh and Retry fields of SOA RR of the primary are set to too high values (in-frequent refresh)
 - frequent zone file change may lead to inconsistency between primary and secondary zone files
- **Zone Thrash**
 - Refresh and Retry fields of SOA RR of the primary are set to too low values (frequent refresh)
 - frequent zone file change may lead to high workload on both primary and secondary and may lead to DoS
- **Information for Targeted Attacks**
 - HINFO and TXT resource records can provide more information for an adversary targeting the organisation
 - e.g. information about software name and version etc.

DNS Data Content: Protective Measures

- analyse DNS RRs data for security implications
- use a zone file integrity checker tool to check for necessary constraints
 - Refresh: secondary server to wait between zone transfers
 - e.g. updated frequently: 20m-2h, infrequently: 2h-12h
 - Retry: secondary server to wait before attempting zone transfer after failure
 - a fraction of Refresh value e.g. 5m-1h
 - Expire: length of time to consider zone information valid if primary server is unreachable
 - a multiple of Refresh e.g. 2w-4w
 - Minimum TTL: the default value of TTL of zone RRsets
 - tells clients how long RRS should be cached
 - depends on how often zone file changes, infrequently (static zone): large e.g. 5d, frequently (dynamic zone): small e.g. 30m
- RRs that leak information should be avoided if possible
- limit information exposure by partitioning of zone files (*split* DNS)
 - two *views*: inside and outside (firewall)
 - drawback: remote users (road warriors) may not be able to use internal servers
- limit information exposure by using separate DNS servers

DNS query/response: Threats

- Forged or Bogus response
 - a compromised authoritative name server
 - compromised host
 - vulnerable DNS software
 - a poisoned cache of a resolving name server
 - spoofed response from attacker with authoritative name server IP address
 - responses from a compromised authoritative name server
 - impact: DoS, client redirection
- Removal of some RRs
 - the adversary removes part of response
 - impact: DoS

DNS query/response: Protective Measures

- underlying problem with threats to DNS query/response is integrity
- to prevent such attacks each response must be verified
 - integrity of the response
 - authenticity of the source
- achieved through DNSSEC
 - authenticate the source: start from a trusted zone such as a root zone and verify the chain down to current zone
 - public key of the trusted zone is *trust anchor*
 - the goal is to verify the public key of the source (of the response)
 - verify the response
 - the response in addition to RRs has an *authenticator*
 - the authenticator is the digital signature of the RRSet of type RRSIG
 - the client verifies the signature using public key of the (response) source
- to protect against removal of RR
 - NSEC or NSEC3 RR is used which lists the RRTypes associated with an owner name and next name in the zone
 - this special RR is also signed and is sent to the resolving name server

DNS Zone Transfer: Protective Measures

- as discussed before is susceptible to DoS
 - limit zone transfer to a set of known entities
 - use of IP address as identity is still vulnerable to IP spoofing
- also vulnerable to tampering
- protected using *transaction signature* or TSIG
 - mutual identification of servers based on a shared secret
 - an adequate method: since number of servers within an administrative domain is limited
 - shared secret is also used to protect/verify the zone transfer transaction
 - a DNSSEC-signed zone only protects RRsets and not all information in a zone file
- another method to protect zone transfer is referred to as SIG(0)
 - similar to RRSIG
 - uses public key cryptography (hence more expensive)
- lower layer protection such as IPsec can also be used

Dynamic Updates: Threats and Countermeasures

Threats

- unauthorised updates
 - adding illegitimate resources to a valid zone
 - deleting legitimate resources
 - altering delegation information
 - NS RR to a child zone
- tampering with the dynamic update content
- replay of captured dynamic updates

Countermeasures

- authentication of entities prevent unauthorised updates and tampering with updates
 - TSIG and SIG(0) mechanisms can provide protection
 - SIG(0) in this case is a better choice
 - including a timestamp in dynamic update limits the replay attack
- use of lower layer protection such as IPsec

References

The materials in this document are reproduced, at times without modification, from the following sources:

- RFC 4880: OpenPGP Message Format
- “S/MIME version 3.2”
- Cryptographic Message Syntax
- RFC 6376: Domain Key Identified Mail (DKIM)
- RFC 4251: SSH Protocol Architecture
- RFC 4253: SSH Transport Layer Protocol
- RFC 4252: SSH Authentication Protocol
- RFC 4254: SSH Connection Protocol
- Bellare, Mihir, Tadayoshi Kohno, and Chanathip Namprempr. “Authenticated encryption in SSH: provably fixing the SSH binary packet protocol.” Proceedings of the 9th ACM conference on Computer and communications security. ACM, 2002.
- RFC 2845: Secret Key Transaction Authentication for DNS (TSIG)

Appendix: SSH Transport Layer Protocol Key Derivation

- Output of Kex is a shared secret K and an exchange hash H
 - H is the hash of the first key exchange
 - H is also used as session identifier (will not change even if keys later change)
- Initialisation vectors:

$IV_{cs} = \text{Hash}(K \parallel H \parallel \text{"A"} \parallel \text{session_id})$

$IV_{sc} = \text{Hash}(K \parallel H \parallel \text{"B"} \parallel \text{session_id})$

- Encryption keys

$KE_{cs} = \text{Hash}(K \parallel H \parallel \text{"C"} \parallel \text{session_id})$

$KE_{sc} = \text{Hash}(K \parallel H \parallel \text{"D"} \parallel \text{session_id})$

- MAC keys

$KI_{cs} = \text{Hash}(K \parallel H \parallel \text{"E"} \parallel \text{session_id})$

$KI_{sc} = \text{Hash}(K \parallel H \parallel \text{"F"} \parallel \text{session_id})$

- if more bits are needed than the hash output then the process is repeated